

CIFAR-10 Dataset

```

In [1]:  import numpy as np
         import matplotlib.pyplot as plt

In [2]:  import keras

In [4]:  from keras.utils import np_utils
         from keras.models import Sequential

In [5]:  from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
         from keras.callbacks import ModelCheckpoint

In [6]:  from keras.datasets import cifar10

In [7]:  (x_train, y_train), (x_test, y_test) = cifar10.load_data() # Need Time

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz (https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz)
170498071/170498071 [=====] - 21s 0us/step

In [8]:  print(x_train.shape, y_train.shape, x_test.shape, y_test.shape)

(50000, 32, 32, 3) (50000, 1) (10000, 32, 32, 3) (10000, 1)

In [9]:  cifar10_labels = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                          'dog', 'frog', 'horse', 'ship', 'truck']

In [10]: y_train[:5]

Out[10]: array([[6],
                [9],
                [9],
                [4],
                [1]], dtype=uint8)

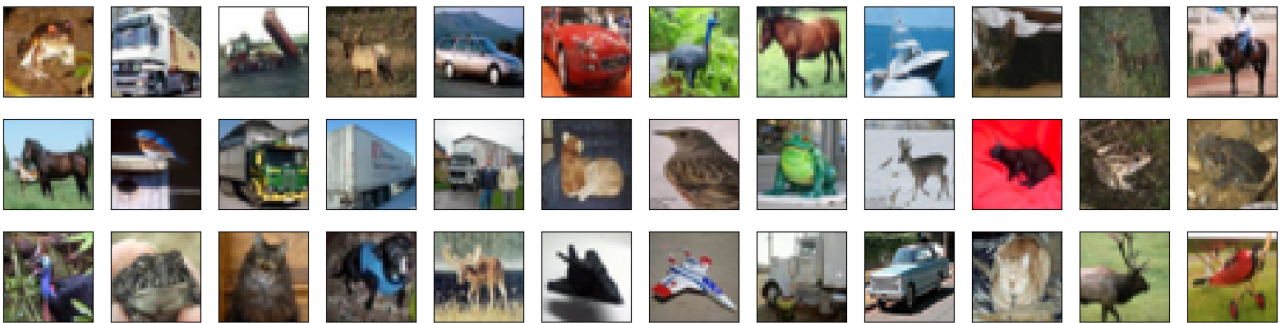
```

2. Show first 24 Training Images

```

In [11]: fig = plt.figure(figsize=(20,5))
         for i in range(36):
             ax = fig.add_subplot(3, 12, i + 1, xticks=[], yticks=[])
             ax.imshow(np.squeeze(x_train[i]))

```



```

In [12]: # rescale [0,255] --> [0,1]
         x_train = x_train.astype('float32')/255
         x_test = x_test.astype('float32')/255

```

4. Split Dataset into Training, Testing, and Validation Sets

```
In [13]: # one-hot encode the labels
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [14]: y_train[:5]
```

```
Out[14]: array([[0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```

```
In [15]: # break training set into training and validation sets
(x_train, x_valid) = x_train[5000:], x_train[:5000]
(y_train, y_valid) = y_train[5000:], y_train[:5000]
```

```
In [16]: print(x_train.shape, y_train.shape, x_valid.shape, y_valid.shape)
```

```
(45000, 32, 32, 3) (45000, 10) (5000, 32, 32, 3) (5000, 10)
```

5. Define the Model Architecture

[\[feature_scaling\]\(cnn-schema1.jpg\)](#)

```
In [17]: from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
```

```
In [18]: model = Sequential()
model.add(Conv2D(filters=16, kernel_size=2,
                  padding='same', activation='relu',
                  input_shape=(32, 32, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2,
                  padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=64, kernel_size=2,
                  padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))

model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(10, activation='softmax'))

#model.summary()
```

In [19]:  model.summary()

Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 16)	208
max_pooling2d (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_1 (Conv2D)	(None, 16, 16, 32)	2080
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_2 (Conv2D)	(None, 8, 8, 64)	8256
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 500)	512500
dropout_1 (Dropout)	(None, 500)	0
dense_1 (Dense)	(None, 10)	5010
Total params: 528,054		
Trainable params: 528,054		
Non-trainable params: 0		

6. Compile the Model

In [20]:  model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])

7. Train the Model

In [21]:  from keras.callbacks import ModelCheckpoint

In [22]:  # train the model
 checkpointer = ModelCheckpoint(filepath='model.weights.best.hdf5', verbose=1, save_best_only=True)
 hist = model.fit(x_train, y_train, batch_size=32, epochs=5,
 validation_data=(x_valid, y_valid), callbacks=[checkerpointer],
 verbose=2, shuffle=True)

Epoch 1/5

Epoch 1: val_loss improved from inf to 1.34982, saving model to model.weights.best.hdf5
 1407/1407 - 20s - loss: 1.6136 - accuracy: 0.4120 - val_loss: 1.3498 - val_accuracy: 0.5258 - 20s/epoch - 14ms/step
 Epoch 2/5

Epoch 2: val_loss improved from 1.34982 to 1.11654, saving model to model.weights.best.hdf5
 1407/1407 - 27s - loss: 1.2820 - accuracy: 0.5410 - val_loss: 1.1165 - val_accuracy: 0.6040 - 27s/epoch - 19ms/step
 Epoch 3/5

Epoch 3: val_loss improved from 1.11654 to 1.04901, saving model to model.weights.best.hdf5
 1407/1407 - 27s - loss: 1.1585 - accuracy: 0.5897 - val_loss: 1.0490 - val_accuracy: 0.6218 - 27s/epoch - 19ms/step
 Epoch 4/5

Epoch 4: val_loss did not improve from 1.04901
 1407/1407 - 26s - loss: 1.0756 - accuracy: 0.6181 - val_loss: 1.1728 - val_accuracy: 0.5936 - 26s/epoch - 18ms/step
 Epoch 5/5

Epoch 5: val_loss improved from 1.04901 to 0.94507, saving model to model.weights.best.hdf5
 1407/1407 - 26s - loss: 1.0311 - accuracy: 0.6369 - val_loss: 0.9451 - val_accuracy: 0.6732 - 26s/epoch - 18ms/step

8. Load the Model with the Best Validation Accuracy

```
In [23]: ▶ # load the weights that yielded the best validation accuracy
model.load_weights('model_weights.best.hdf5')
```

9. Test Accuracy rate

```
In [24]: ▶ # evaluate and print test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('\n', 'Test accuracy:', score[1])
```

Test accuracy: 0.6646000146865845