

ISE 599 Deep Learning Student ID: 7636428840

## Diabetes dataset

```
In [69]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [70]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [71]: from tensorflow import keras
```

```
In [72]: from tensorflow.keras import layers
```

```
In [73]: # see keras version
keras.__version__
```

```
Out[73]: '2.12.0'
```

```
In [74]: #pd.set_option('display.max_columns', None)
```

### Get data

```
In [75]: df = pd.read_csv('diabetes.csv')
df
```

```
Out[75]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	PedigreeFunc	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

The diabetes.csv comes from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements. All patients here are females at least 21 years old of Pima India heritage. The variables are:

Pregnancies: the number of times pregnant

Glucose Plasma: glucose concentration at 2 hours in an oral glucose tolerance test.

BloodPressure: Diastolic blood pressure (mm Hg)

SkinThickness: Triceps skin fold thickness (mm)

Insulin: 2-Hour serum insulin (mu U/ml)

BMI: Body mass index(weight in kg/(height in m)^2)

DiabetesPedigreeFunction

Age

Outcome(0 or 1)

In [76]:  df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Pregnancies      768 non-null    int64
1   Glucose          768 non-null    int64
2   BloodPressure    768 non-null    int64
3   SkinThickness    768 non-null    int64
4   Insulin          768 non-null    int64
5   BMI              768 non-null    float64
6   PedigreeFunc     768 non-null    float64
7   Age              768 non-null    int64
8   Outcome          768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [77]:  df.isnull().sum() # It has no missing values, because all missing values are replaced by 0.

```
Out[77]: Pregnancies      0
         Glucose          0
         BloodPressure    0
         SkinThickness    0
         Insulin          0
         BMI              0
         PedigreeFunc     0
         Age              0
         Outcome          0
         dtype: int64
```

In [78]:  # How many diabetes patients?In [79]:  df.Outcome.value\_counts()


```
Out[79]: 0    500
         1    268
         Name: Outcome, dtype: int64
```


In [80]:  df.Outcome.value\_counts()/768

```
Out[80]: 0    0.651042
         1    0.348958
         Name: Outcome, dtype: float64
```

In [81]:  # There are 35% diabetes patients

## Data Preparation

```
In [82]:  # looking for missing values and outliers)
         # Some columns have entries equal to zero
         # For some columns that is not possible (i.e., BMI, Insulin)
```

```
In [83]:  for col in df.columns:
         zeros = df.loc[df[col]==0].shape[0]
         print(col+" : "+str(zeros))
```

```
Pregnancies: 111
Glucose: 5
BloodPressure: 35
SkinThickness: 227
Insulin: 374
BMI: 11
PedigreeFunc: 0
Age: 0
Outcome: 500
```

In [84]:  # Imputation

```
In [85]: # replace the zeros with nan
```

```
In [86]: df['Glucose'] = df['Glucose'].replace(0, np.nan)
df['BloodPressure'] = df['BloodPressure'].replace(0, np.nan)
df['SkinThickness'] = df['SkinThickness'].replace(0, np.nan)
df['Insulin'] = df['Insulin'].replace(0, np.nan)
df['BMI'] = df['BMI'].replace(0, np.nan)
```

```
In [87]: df.isnull().sum()
```

```
Out[87]: Pregnancies      0
Glucose      5
BloodPressure  35
SkinThickness 227
Insulin      374
BMI          11
PedigreeFunc   0
Age           0
Outcome       0
dtype: int64
```

```
In [88]: # replace the nan with the average of that column
```

```
In [89]: df['Glucose'] = df['Glucose'].fillna(df['Glucose'].mean())
df['BloodPressure'] = df['BloodPressure'].fillna(df['BloodPressure'].mean())
df['SkinThickness'] = df['SkinThickness'].fillna(df['SkinThickness'].mean())
df['Insulin'] = df['Insulin'].fillna(df['Insulin'].mean())
df['BMI'] = df['BMI'].fillna(df['BMI'].mean())
```

## Split data

```
In [90]: X = df.drop(['Outcome'], axis = 1)
Y = df.Outcome
```

```
In [91]: X_train, X_test, y_train, y_test = train_test_split(X, Y, stratify = Y,
test_size=0.2,
random_state=1)
```

```
In [92]: y_train.shape
```

```
Out[92]: (614,)
```

```
In [93]: y_test.shape
```

```
Out[93]: (154,)
```

## Scale the data

```
In [94]: scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [ ]:
```

## Build Model

Build the Dense Neural Network (DNN) with 2 hidden layers

```
In [95]: network1 = keras.Sequential([
layers.Dense(32, activation='relu'), # first hidden layer
layers.Dense(16, activation='relu'), # second hidden layer
layers.Dense(1, activation='sigmoid') # output layer
])
```

**Compile Model**

```
In [96]: # use loss function binary_crossentropy
network1.compile(optimizer='rmsprop',
                 loss='binary_crossentropy',
                 metrics=['accuracy'])
```

**train(fit) model**

```
In [97]: n_epochs = 55
```


```
In [98]: history = network1.fit(X_train_scaled, y_train,
                               epochs=n_epochs, batch_size=20,
                               validation_split = 0.20,
                               verbose = 0,
                               )
```

```
In [99]: # values for train loss and train accuracy are shown at each step
```

```
In [100]: loss = history.history["loss"]
val_loss = history.history["val_loss"]
```

```
In [101]: epochs = range(1, n_epochs+1)
```

```
In [102]: plt.figure(figsize=(10,5))
plt.plot(epochs, loss, "r",
        label="Training loss")
plt.plot(epochs, val_loss, "b",
        label="Validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid()
```



```
In [103]: # The model starts overfitting after x epochs
```

```
In [104]: ## Retrain model with all train data (with x epochs)
```

```
In [105]: model = keras.Sequential([
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

y_pred_probabilities = model.fit(X_train_scaled, y_train, epochs=10, batch_size=20)
print(y_pred_probabilities)
```

Epoch 1/10  
 31/31 [=====] - 1s 3ms/step - loss: 0.6411 - accuracy: 0.6336  
 Epoch 2/10  
 31/31 [=====] - 0s 4ms/step - loss: 0.5414 - accuracy: 0.7443  
 Epoch 3/10  
 31/31 [=====] - 0s 3ms/step - loss: 0.4992 - accuracy: 0.7638  
 Epoch 4/10  
 31/31 [=====] - 0s 4ms/step - loss: 0.4763 - accuracy: 0.7687  
 Epoch 5/10  
 31/31 [=====] - 0s 2ms/step - loss: 0.4647 - accuracy: 0.7785  
 Epoch 6/10  
 31/31 [=====] - 0s 2ms/step - loss: 0.4561 - accuracy: 0.7801  
 Epoch 7/10  
 31/31 [=====] - 0s 4ms/step - loss: 0.4504 - accuracy: 0.7818  
 Epoch 8/10  
 31/31 [=====] - 0s 3ms/step - loss: 0.4455 - accuracy: 0.7915  
 Epoch 9/10  
 31/31 [=====] - 0s 4ms/step - loss: 0.4415 - accuracy: 0.7932  
 Epoch 10/10  
 31/31 [=====] - 0s 4ms/step - loss: 0.4380 - accuracy: 0.7948  
 <keras.callbacks.History object at 0x000002438F083A00>

```
In [106]: import pandas as pd
yhat = model.predict(X_test_scaled)

y_test = np.array(y_test)
yhat = np.array(yhat)

df2 = pd.DataFrame({'y_test': y_test, 'yhat': yhat})
df2['y_test'] = pd.Series(y_test)
df2['yhat'] = pd.Series(yhat)
print(df2)
```

5/5 [=====] - 0s 3ms/step

	y_test	yhat
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
..	...	...
149	0	0
150	0	0
151	0	0
152	1	1
153	0	0

[154 rows x 2 columns]

```
In [107]: confusion_matrix = pd.crosstab(y_test, yhat, rownames=['Actual'], colnames=['Predicted'])
print(confusion_matrix)
```


	Predicted 0	Predicted 1
Actual 0	100	0
Actual 1	0	54

```
In [108]: # 2. Add a column with the error rates (accuracy) for each type of patient:
confusion_matrix["Total"] = confusion_matrix.sum(axis=1)
confusion_matrix["Error Rate"] = 1 - confusion_matrix[1] / confusion_matrix["Total"]
print(confusion_matrix)
```

	Predicted 0	Predicted 1	Total	Error Rate
Actual 0	100	0	100	1.0
Actual 1	0	54	54	0.0

```
In [109]: import seaborn as sns
from sklearn.metrics import confusion_matrix
mat = confusion_matrix(y_test, yhat)
plt.figure(figsize=(7, 5))
sns.heatmap(mat, annot=True)
```

Out[109]: <Axes: >



what is the accuracy rate for predicting if a patient has diabetes, if in fact he has diabetes?

The accuracy rate for incorrection is 0, so there are no cases of incorrectly predict the diabetes status.

```
In [110]: from sklearn.metrics import classification_report
target_names = ['Diabetes', 'Normal']
print(classification_report(y_test, yhat, target_names=target_names))
```


	precision	recall	f1-score	support
Diabetes	1.00	1.00	1.00	100
Normal	1.00	1.00	1.00	54
accuracy			1.00	154
macro avg	1.00	1.00	1.00	154
weighted avg	1.00	1.00	1.00	154

```
In [111]: from sklearn.metrics import roc_curve

y_pred_keras = model.predict(X_test_scaled).ravel()
fpr, tpr, thresholds = roc_curve(y_test, y_pred_keras)

plt.plot([0,1], [0,1], 'k-')
plt.plot(fpr, tpr, label='Knn')
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('ROC curve')
plt.grid()
plt.show()
```

5/5 [=====] - 0s 2ms/step



```
In [112]: from sklearn.metrics import roc_auc_score
roc_auc_score(y_test, y_pred_keras)
```

Out[112]: 0.827962962962963

```
In [113]: # define a function that accepts a threshold and prints sensitivity and specificity
def evaluate_threshold(threshold):
    print('Sensitivity:', tpr[thresholds > threshold][-1])
    print('Specificity:', 1 - fpr[thresholds > threshold][-1])

evaluate_threshold(0.3)
```

Sensitivity: 0.7962962962962963  
Specificity: 0.72

```
In [114]: evaluate_threshold(0.35)
```

Sensitivity: 0.7222222222222222  
Specificity: 0.8

```
In [115]: evaluate_threshold(0.4)
```

Sensitivity: 0.5925925925925926  
Specificity: 0.83

```
In [116]: evaluate_threshold(0.5)
```

Sensitivity: 0.48148148148148145  
Specificity: 0.86

**test model**

```
In [117]: test_loss, test_acc = model.evaluate(X_test_scaled, y_test)
```

```
5/5 [=====] - 0s 4ms/step - loss: 0.5033 - accuracy: 0.7273
```

```
In [118]: test_acc # 0.7597402334213257
```

```
Out[118]: 0.7272727489471436
```

```
In [119]: test_loss # 0.5030236840248108
```

```
Out[119]: 0.5032797455787659
```